



Model Training Health Diagnostic — final_model_20_evvp_18k.pt

Architecture: 7→128d, 3 layers, 487,553 params (enc 396,928 / dec 82,304 / val 8,321)

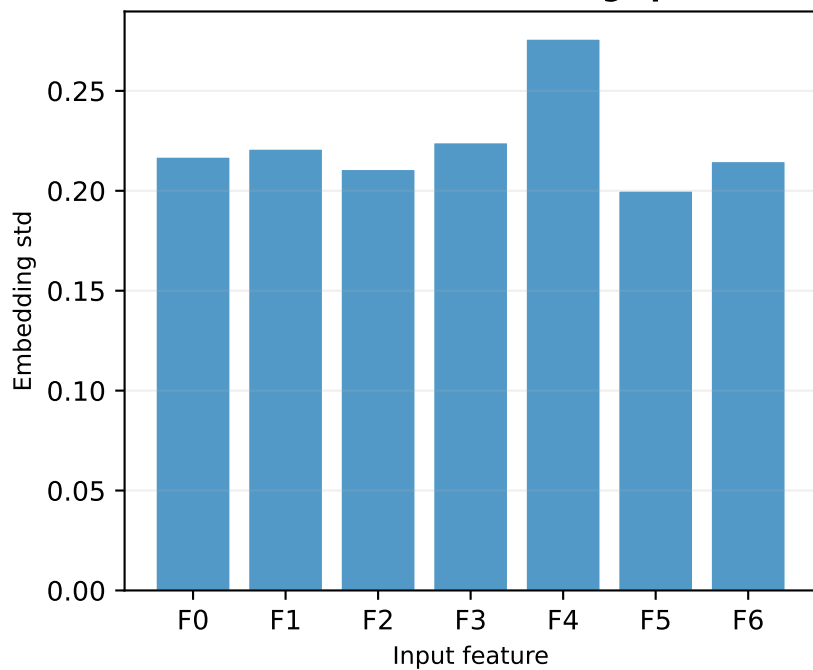
A. Health scores



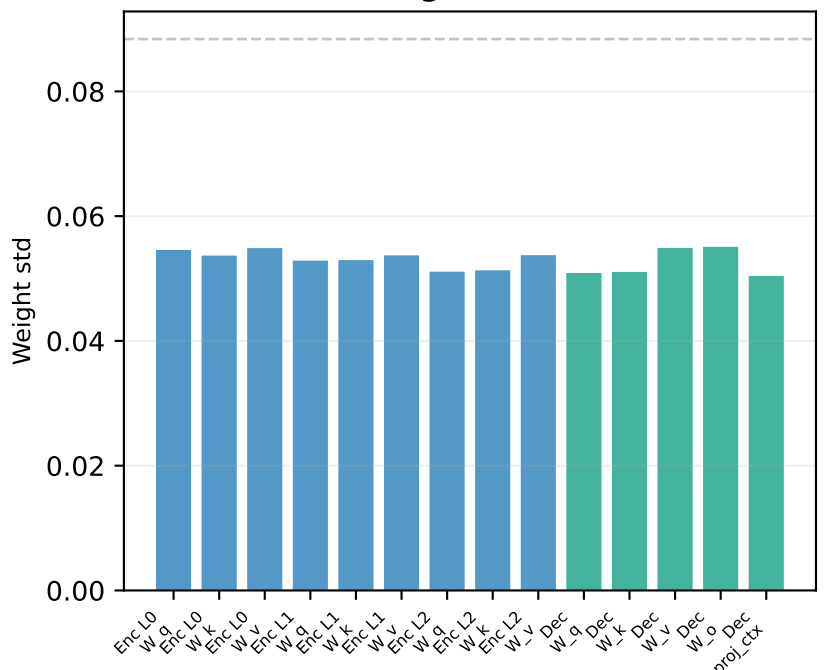
B. Attention weight std per layer



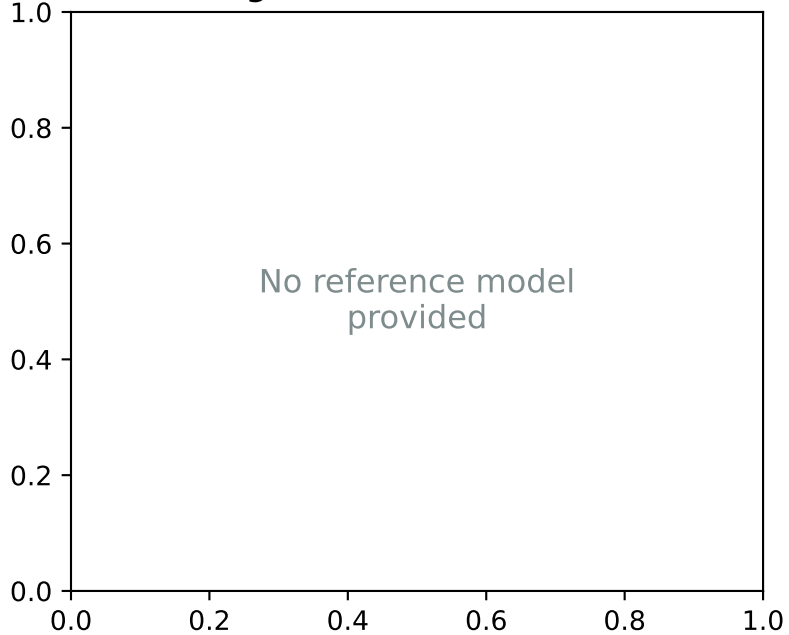
C. Per-feature embedding spread



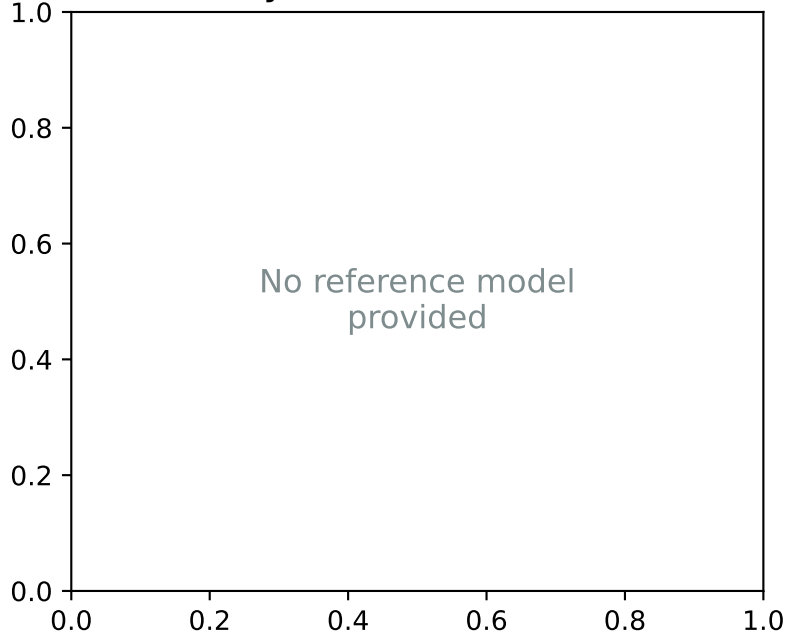
D. All attention weights: encoder vs decoder



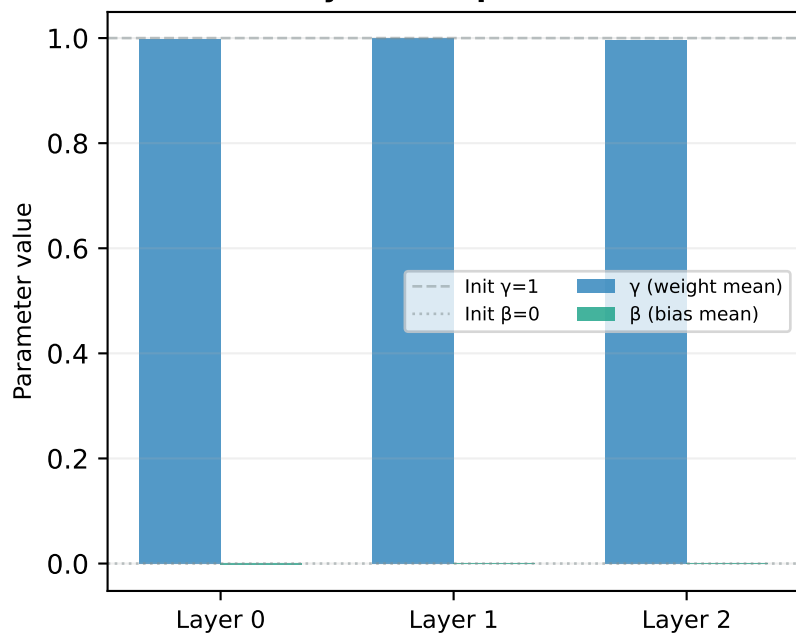
E. Weight drift (needs reference)



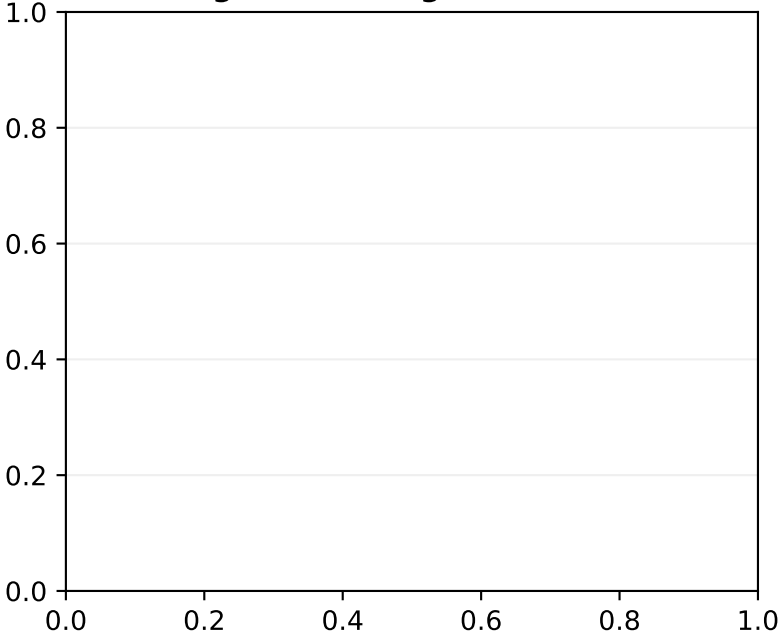
F. Per-layer drift (needs reference)



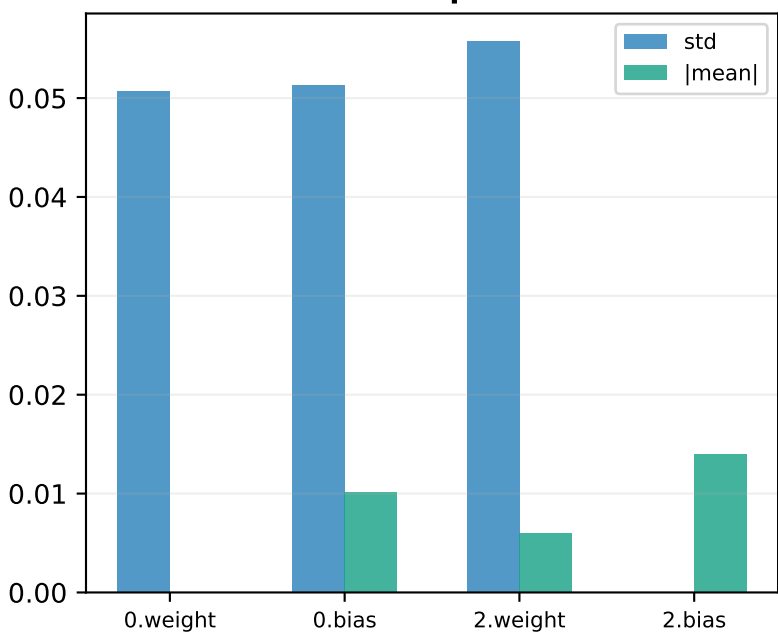
G. LayerNorm parameters



H. Weight std change from reference



I. Value head parameters



Overall Training Readiness: 57/100

DEVELOPING

- Encoder layers have NOT specialized — attention weights are near-identical across layers
- Decoder has moved well beyond initialization ✓
- Input features show differentiated embeddings ✓

RECOMMENDATION: Continue training. Model is developing but not yet converged.